

Multi-Tenant SaaS Dashboard

— Proposal Request

We're selecting a development partner to build a multi-tenant marketing intelligence dashboard, and we'd like your proposal — approach, team, timeline and cost — based on the brief below.

Project Brief

We're building a unified marketing intelligence platform — a single dashboard that pulls marketing data from Meta, Instagram, LinkedIn and Google Analytics into one place, replacing the manual work of collecting and reconciling reports across platforms today.

Problem: Marketing data is scattered across social, ad and analytics platforms, each with its own terminology and metrics. Teams manually download and consolidate reports, historical comparison is hard, and reporting depends on whoever happens to hold the knowledge.

Target users:

- **Marketing Manager** — campaign and channel performance
- **Marketing Head / CMO** — high-level overview, channel comparison, ROI, executive reporting
- **Business Owner / Founder** — a simple view of leads, traffic and engagement
- **Agencies** — managing multiple clients from one account, with client-specific dashboards

Product principles: simple (no API/analytics expertise required to use it), unified (one place for everything), historical (a real stored data layer, not just live API pulls on load), automated (minimal manual intervention), and — eventually — actionable (AI-driven explanations and recommendations, later phase).

Requirements Summary

Phase 1 scope — data integration and unified dashboard

Connect and pull metrics from:

- **Meta/Facebook** — followers, reach, impressions, engagement, post-level performance
- **Instagram** — followers, reach, impressions, profile visits, engagement, post performance

- **LinkedIn** — followers, impressions, clicks, reactions, shares, post performance
- **Google Analytics 4** — users, sessions, engagement rate, traffic sources, landing pages, device data

Architecture: No live API calls on every dashboard load. Data flows through connectors → normalization → database → dashboard API → UI, so the dashboard reads from stored historical data, not the source platforms directly.

Sync: Centralized scheduler, configurable frequency (1–24 hours, default 6), plus a manual "sync now."
 Sync process: check token status → refresh if supported → call platform API → validate → normalize → store → log sync status.

Auth & tokens: OAuth per platform where supported. Tokens encrypted at rest, auto-refreshed where the platform allows it, with a clear "reconnect" path when re-authentication is needed — no unexplained API errors surfaced to the user.

Core data model: Organizations, Users (with roles), Connected Accounts, Daily Metrics, Content/Posts, Sync Logs.

User roles: Admin (manage org, connections, users, full data access), Manager (view + trigger sync), Viewer (view only).

Multi-tenancy: Built as a multi-tenant SaaS from day one. Every data object is scoped Organization → Account → Platform → Data. One customer must never be able to access another customer's data — this is a hard requirement, not a later hardening pass.

Natural-language insights: Beyond the dashboard, users need to be able to ask plain-language questions about their own marketing performance — e.g. "which channel drove the most engagement last month?" or "why did LinkedIn impressions drop in March?" — and get an answer grounded in the stored structured data. This is a defining part of the product, not a nice-to-have tacked on at the end, and we want your realistic read on it now: what it takes to build and maintain reliably, what it costs, and how it fits your team's current expertise. Implementation can be sequenced after the core dashboard ships, but we're evaluating proposals on how seriously this part is addressed, not just the dashboard build.

Build priority order: (1) multi-tenant foundation, auth, org management, OAuth/token framework — (2) the four platform integrations — (3) data pipeline: connectors, normalization, scheduler, error handling — (4) dashboard UI — (5) reliability: monitoring, sync logs, retry, connection health — (6) natural-language insights layer (see above), anomaly detection.

Out of scope for Phase 1: ad platform data (Meta/Google/LinkedIn Ads), CRM/lead system integration, YouTube, Search Console — these are Phase 2.

What We're Looking For

A partner who can own this end-to-end — architecture through delivery — and who has actually built multi-tenant SaaS products before, not just single-tenant client sites. We're evaluating on:

- Relevant past work (ideally something with OAuth integrations, a sync/scheduler pattern, and real multi-tenancy — not just role-based access on a shared database)
- How you'd approach the architecture given the constraints above
- Realistic timeline and cost, broken down by phase
- How you handle ongoing maintenance and change requests after launch, not just the initial build

We're not looking for the cheapest bid — we're looking for the team that's actually built something like this before and can defend their approach.

Proposal Questions

Please answer these alongside your proposal:

1. **Similar work** — Share 1–2 projects most similar to this (multi-tenant SaaS, OAuth integrations, or a reporting/analytics dashboard). What was genuinely hard about them?
2. **Tech stack** — What would you build this on (backend, database, frontend), and why that stack for this specific brief?
3. **Multi-tenancy approach** — How would you implement tenant isolation concretely? Row-level policies, per-tenant schemas, an app-layer scoping rule — and what's your experience with it holding up in production?
4. **Team & structure** — Who works on this (roles, seniority), and is it a dedicated team or shared across other client work?
5. **Timeline** — Realistic delivery estimate by phase (foundation → integrations → pipeline → dashboard → reliability), not just a single end date.
6. **Cost** — Implementation cost broken down by phase, and your model for it (fixed bid, time & materials, retainer).
7. **Maintenance** — What ongoing maintenance/support looks like after launch, and

how change requests (e.g., adding a field, adding a platform) are typically scoped and priced.

8. **AI-assisted delivery** — What AI coding tools or assistants does your team use in day-to-day delivery, and how has that changed your typical timeline or team size over the last year?
9. **Backend platform experience** — Have you built on a schema-driven or generated-backend platform (e.g., Supabase, Firebase, Xano, or similar) versus a fully hand-rolled backend? What's your honest take on the tradeoffs — speed vs. control — for a project like this?
10. **Infra & hosting** — Which cloud or backend platform would you use for this, and roughly what would it cost per tenant to run the infrastructure at moderate scale — excluding AI token costs?
11. **AI/agent build** — Do you have in-house expertise building AI agents? What tech stack and team expertise would you bring to the natural-language insights layer described above — where users ask questions in plain language and get answers grounded in the structured marketing data? Roughly how much engineering effort (time, team, and any tooling like MCP) would that layer take relative to the rest of the build?
12. **Risks** — What would you flag as the riskiest or least-defined part of this brief as written?

Reply with your proposal and answers to the questions above, and call out any assumptions you're making where the brief doesn't fully specify something.

We've received 500+ initial responses to our outreach, so we'd prefer to keep communication to email for now while we shortlist down to a handful of agencies to talk with directly. This process will likely take a few weeks to a couple of months — thanks in advance for your patience.